

ThrivikramaRao Kavuri

thrivikramaraokavuri@gmail.com | 716-253-5373 | Thrivikramarao Kavuri | Vikram Kavuri | TR



Technical Skills

- **Data & Cloud Platforms:** Azure, AWS (S3, Glue), Snowflake, PostgreSQL, Microsoft SQL Server
- **ETL Tools:** Azure Data Factory, Microsoft Fabric, PySpark, Apache Airflow, Databricks, SSIS
- **BI Tools:** Tableau, Power BI, Qlik Sense, SSRS, Apache Superset, Microsoft Excel
- **Data Engineering Skills:** Data Migration, HIPAA, GDPR, Scalable ETL pipeline design, Data Quality (QA/QC)

Experience

Data Analytics Engineer | The Arc Erie County New York — New York, US Jan 2025 – Present

- Engineered **Airflow** pipelines replacing manual extracts, pulled PrecisionCare and NetSuite **REST APIs** into **Snowflake**, keeping 500K rows moving with retries and monitoring, and mapped outputs for **Microsoft Fabric**.
- Designed **Snowflake** dimensional models and governed views for billing analytics, then drafted a Microsoft Fabric Lakehouse migration plan, aligning schemas, RBAC, and masking so teams could move safely.
- Validated lakehouse tables with **Pandas** rules and **SQL** tests when quality drifted built scorecards with alerts and fixed root causes early so downstream reporting stayed accurate and audit-ready.
- Translated billing workflows into secure **Snowflake** views and **Tableau** dashboards, coached users on KPI definitions, and documented pipelines and lineage so stakeholders could self-serve answers without rework.

Data Science Analyst | Accenture India Pvt Ltd — Bangalore, India May 2021 – Jan 2023

- Coordinated on-prem **SQL** data warehouse migration to **Azure** using **Azure DevOps** pipelines, scripted schema and data moves, ran validation queries, and coordinated cutover so BI teams could switch with minimal downtime.
- Engineered **PySpark** ELT in **AWS Glue** when daily loads kept failing, added incremental pulls, retry logic, and schema checks so **S3** to **Redshift** and **Snowflake** feeds stayed reliable for each run and migration ready.
- Designed **Snowflake** star models and **SQL** views after exec metrics disagreed, standardized KPI logic, tuned clustering and caching, and brought report refresh time down to under an hour without breaking totals.
- Migrated **QlikSense** reporting to **Power BI**, partnered with stakeholders to clarify KPI definitions, documented measures and data rules, and rebuilt datasets so dashboards matched results and stayed trusted.

Business Analyst | SRIT Pvt Ltd — Bangalore, India Jan 2020 – Apr 2021

- Translated clinical and ops needs into BRDs and FRDs with **SQL**-ready KPIs and acceptance criteria, enabling engineers to build ETL and reporting with fewer reworks while keeping scope clear for stakeholders.
- Developed **Power BI** and **Apache Superset** dashboards over 1M+ healthcare records, tuning queries and visuals so executives saw near-real-time patient care KPIs daily and response time improved 8%.
- Organized stakeholder workshops to reconcile conflicting metric definitions, documenting lineage and refresh schedules so analysts had consistent datasets and 1M+ record dashboards stayed trusted.

Education

Masters in Data Science, University at Buffalo, GPA: 3.7/4.0 Jan 2023 – May 2024

Bachelor's in Computer Science, Karunya University, GPA: 3.4/4.0 Jun 2016 – Jul 2020

Projects

Telco Customer Churn Prediction (Tools: Python, Flask, Random Forest, XGBoost, Excel)



- Established a churn prediction workflow using Flask and XGBoost (87% ROC-AUC), modeling customer contracts, add-ons, and payments to identify high-risk cohorts and recommend targeted retention actions.

Real-Time Stock Market Data Pipeline (Tools: PySpark, Databricks, GCP, Apache Airflow)



- Devised a real-time stock market data pipeline, leveraging PySpark for distributed processing, GCP Big Query for scalable storage, Databricks for optimized querying, and Airflow for orchestration.

Certifications

- Generative AI Professional (Oracle)
- Fabric Data Engineer Associate (Azure)